

EngSci 721

Inverse Problems and Learning From Data

Oliver Maclaren (oliver.maclaren@auckland.ac.nz)

1. Basic concepts [5 lectures + 1 Tutorial]

Forward vs inverse problems. Well-posed vs ill-posed problems. Algebra and calculus of inverse problems (left and right inverses, generalised and pseudo inverses, resolution operators, matrix calculus). Representing higher dimensional problems (image data etc).

2. Instability and regularisation in linear and nonlinear problems [6 lectures + 1 Tutorial]

Instability and related issues for generalised inverses. Introduction to regularisation and trade-offs. Tikhonov regularisation. Higher-order Tikhonov regularisation. Sparsity and regularisation using different norms. Truncated singular value decomposition. Iterative regularisation, including stochastic/mini-batch gradient descent.

3. Further topics [3 lectures + 1 Tutorial]

Regularisation parameter choice, including statistical and machine learning views of regularisation. Confidence sets for linear and nonlinear models. Physics-informed machine learning and neural networks.

Module overview

Inverse Problems and Learning From Data (*Oliver Maclaren*)

[~14 lectures/3 tutorials]

Lecture 13: Nonlinear models

Topics:

- Fitting a simple nonlinear scientific model (ODE)
- Constructing confidence sets
- Complications

EngSci 721 : Lecture 13 : Nonlinear models

Fitting simple nonlinear models to data

Eg 'How can I fit a nonlinear ODE to data?'

'What about stochastic/
large models?'

→ simple practical example

→ pointers to lit. for more complex cases

- └ derivative calculation
- └ surrogates/emulators
- └ ensemble methods
- └ scientific machine learning (L14)
- + Ru's section (Bayes)

Nonlinear?

We have indicated several times that our ideas apply to both linear & nonlinear (in parameters) models

→ however we have spent a lot of time on linear models!

Let's look at a 'simple' nonlinear example, where number of params < data points, ie the nonlinear version of a 'tall' system

We will also consider some elementary statistical ideas like noise & confidence intervals

Logistic growth



'S shaped'
growth curve

S-shaped growth curves describe
many physical & biological
phenomena, e.g.

- increase of cell number in cell culture
 - population growth in ecology
 - infection curves in disease dynamics
 - reaction rate in enzyme kinetics
- etc!

The simplest (but not only!) such model is given by the logistic growth model:

$$\frac{dc}{dt} = r c \left(1 - \frac{c}{K}\right)$$

$$c(0) = c_0$$

c = count or concentration

r = growth rate

K = 'carrying' capacity (saturation value)

Logistic growth.

Technically we can solve this
analytically for:

$$c(t) = \frac{K c_0}{c_0 + (K - c_0) e^{-rt}}$$

Notes:

• $t \rightarrow \infty, c(t) \rightarrow K$

• c/K small \Rightarrow

$$\frac{dc}{dt} = r c \left(1 - \frac{c}{K}\right) \approx r c$$

• model is nonlinear in parameters

$$r, c$$

• data have form $(t, c)_{i=1, \dots, n}$

However, pretend we don't know
analytical soln!

Trajectory matching

A simple way to fit this model is by nonlinear least squares:

$$\min_{\theta} \|y^{\text{data}} - y^{\text{model}}(\theta)\|^2 \quad \left. \right\} \text{nonlinear in } \theta$$

where . $\theta = (r, K)$

$$y^{\text{data}} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_d \end{bmatrix} = \begin{bmatrix} y(t_1) \\ y(t_2) \\ \vdots \\ y(t_d) \end{bmatrix}$$

$$y^{\text{model}}(\theta) = \begin{bmatrix} y^{\text{model}}(t_1, \theta) \\ \vdots \\ y^{\text{model}}(t_d, \theta) \end{bmatrix}$$

→ If we don't know $y^{\text{model}}(\theta)$ analytically we simply call a numerical solver for the ODE to evaluate for any parameter

→ The above is sometimes called 'trajectory matching', as we aim to fit the 'trajectory' to data.

Optimisation

For small / fast to solve problems we can just use standard libraries to solve.

→ nonlinear least squares form:
can use 'least-squares' function from `scipy.optimize` (python). Similar functions exist in matlab, Julia, R etc.

→ for more general objective functions can use eg 'minimize' from `scipy.optimize` (python) or eg `fminsearch` (matlab) etc.

→ these interface to various derivative-based & derivative-free optimisers.

↳ will compute derivatives for you if needed.

Derivative calculator?

For standard languages & optimizers, if you need derivatives with respect to a function $f(y(\theta))$ of the model output $y(\theta)$, the library will typically compute them using finite differences, e.g. for step-size h ,

$$\frac{\partial f(y(\theta))}{\partial \theta_i} \approx \frac{f(y(\theta + h\epsilon_i)) - f(y(\theta))}{h}$$

where ϵ_i is the unit vector in the i th dir.

- If f is vector-valued, this gives a column of derivatives for each θ_i

$$\frac{\partial f}{\partial \theta^T} = \begin{bmatrix} f_1 \\ f_2 \\ f_3 \end{bmatrix} \left[\begin{array}{c} \theta_1 \\ \theta_2 \\ \vdots \end{array} \right]$$

- If scalar-valued it will just be a number for each θ_i :

$$\frac{\partial f}{\partial \theta^T} = f \left[\begin{array}{c} \theta_1 \\ \theta_2 \\ \vdots \end{array} \right]$$

Derivative calculator

This finite difference approach

- has a cost that scales with number of parameters, for each iteration

↳ OK for small, cheap problems, expensive for large & or costly to run models

- noisy/inaccurate & dependent on choice of h

More sophisticated approaches:

- automatic/algorithmsic
- symbolic/analytical

Modern ML libraries like tensorflow / pytorch etc are really 'just' autodiff libraries!

+ Julia has 'native' autodiff !

Idea: automatically differentiate any code you write !

Example

```
# logistic
from scipy.integrate import solve_ivp

def dydt_p(t,y,prm):
    # params
    r = prm[0]*1E-3 #scale relative to 1E-3
    K = prm[1]*10
    # vars
    C = y[0]
    # ode
    dydt = r*C*(1-C/K)
    return dydt
```

```
# ODE for given param.
dydt = lambda t, y: dydt_p(t,y,prm)
✓ 0.0s
```

```
# parameters
r_true = 2.5
K_true = 8
prm_true = np.array([r_true,K_true])

# initial condition
C0 = 0.7
y0 = [C0]

# ODE for given param. (closure...)
dydt = lambda t, y: dydt_p(t,y,prm_true)

# t_data
t_span=[0,4000]
t_grid = np.arange(t_span[0],t_span[1],1)
t_data = np.arange(0,4000,500)

# true solution for given param
solution_object = solve_ivp(fun=dydt,t_span=t_span,t_eval=t_data,y0=y0,dense_output=True)
sol = solution_object.sol

sigma = 5
y_true = sol(t_grid)[0]
y_data = sol(t_data)[0] + np.random.normal(loc=0,scale=sigma,size=len(t_data))
```

example

```
# best fit
def residual(prm):
    # ODE for given param. (closure...)
    dydt = lambda t, y: dydt_p(t,y,prm)

    # t_data
    t_span=[0,4000]
    t_data = np.arange(0,4000,500)

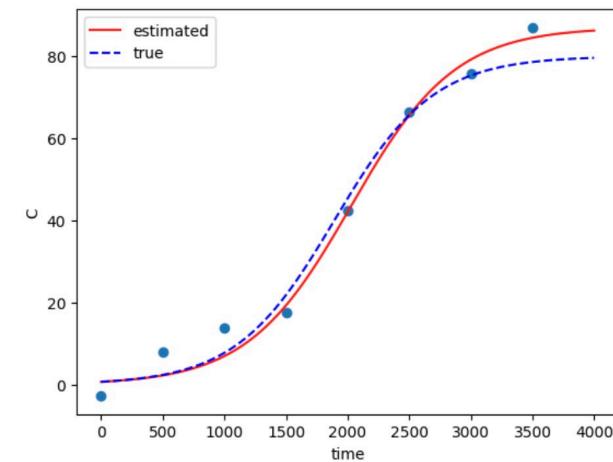
    # true solution for given param
    solution_object = solve_ivp(fun=dydt,t_span=t_span,t_eval=t_data,y0=y0,dense_output=True)
    sol = solution_object.sol

    y_pred = sol(t_data)[0]

    return y_data - y_pred
```

```
# nonlinear least squares
from scipy.optimize import least_squares
x_initial = np.array([1,1])
result = least_squares(fun=residual,x0=x_initial)
prm_est = result.x

# ODE for given param. (closure...)
dydt = lambda t, y: dydt_p(t,y,prm_est)
# est solution for given param
solution_object = solve_ivp(fun=dydt,t_span=t_span,t_eval=t_data,y0=y0,dense_output=True)
sol = solution_object.sol
plt.plot(t_data,y_data,'o')
plt.plot(t_grid,sol(t_grid)[0],'r-',label='estimated')
plt.plot(t_grid,y_true,'b--',label='true')
plt.xlabel('time')
plt.ylabel('C')
print(prm_est)
```



Confidence sets

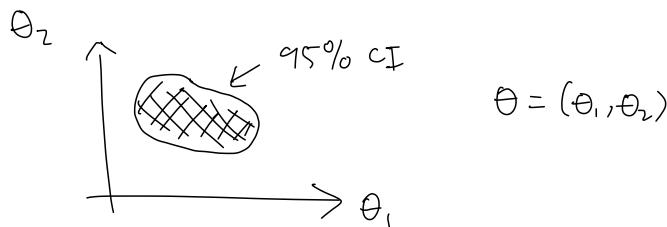
As discussed in the previous lecture,

If we know the noise level, we can form ' $1-\alpha$ ' confidence sets (eg $95=1-\delta$) by 'hypothesis test inversion':

$$CI(\theta) = \left\{ \theta : \frac{1}{2} \|y^{obs} - y(\theta)\|^2 \leq c_{1-\alpha} \right\}$$

for 'cut-off' $c_{1-\alpha}$ given by the $1-\alpha$ threshold of the chi-square distribution

(e.g. confidence set = all non-rejected parameters)



→ Applies to nonlinear models with additive error:

$$y = f(\theta) + \epsilon$$

→ can generalise using likelihood functions (see R)

Example

Here we can just 'grid up':

```
r_min = 0
r_max = 10
r_grid = np.linspace(r_min,r_max,500)

K_min = 5
K_max = 15
K_grid = np.linspace(K_min,K_max,500)

nllikes = np.zeros((len(r_grid),len(K_grid)))
for i, ai in enumerate(r_grid):
    for j, bj in enumerate(K_grid):
        nllikes[i,j] = nllikelihood((ai,bj))

✓ 4m 42.0s
```

(would have to sample or profile in high dim)

Plot contour of 95% boundary:

```
from scipy.stats import chi2
delta = chi2.ppf(q=0.95, df=len(y_data))

contours = plt.contour(K_grid,r_grid, nllikes, levels=[delta], colors='black')

plt.xlabel(r'$\theta_2$', fontsize=16)
plt.ylabel(r'$\theta_1$', fontsize=16)

plt.plot(prm_est[1],prm_est[0],marker='x',color='red',label='estimate')
plt.plot(prm_true[1],prm_true[0],marker='o',color='blue',label='true')
plt.legend()
plt.show()
```



Example

```

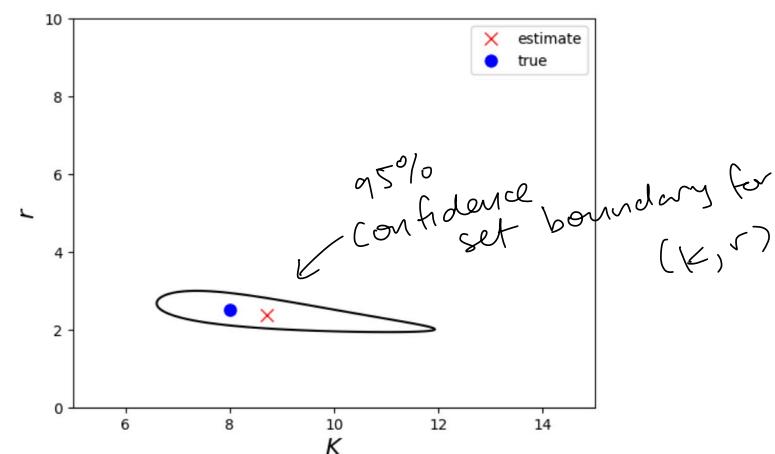
from scipy.stats import chi2
delta = chi2.ppf(q=0.95, df=len(y_data))

contours = plt.contour(K_grid, r_grid, nllikes, levels=[delta], colors='black')

plt.xlabel(r'$K$', fontsize=16)
plt.ylabel(r'$r$', fontsize=16)

plt.plot(prm_est[1], prm_est[0], markersize=8, color='r', marker='x', linestyle='None', label='estimate')
plt.plot(prm_true[1], prm_true[0], markersize=8, color='b', marker='o', linestyle='None', label='true')
plt.legend()
plt.show()

```

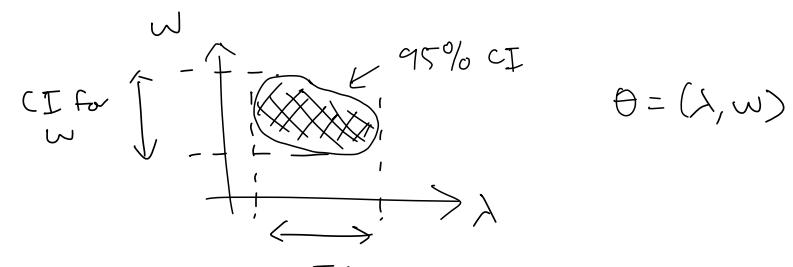


Note: choose units for r & K to make similar scale
 $(K$ actually 80, we use $\bar{K}=8$ & $K=10\bar{K})$

Profile projections for 'targeted CIs'

we can 'project' the full parameter confidence set onto the individual parameter axes to get conservative ($\geq 1-\alpha$ coverage) CIs for individual param.

Idea: We can't reject a subcomponent (parameter) λ of a parameter vector θ if there exists a 'complementary' parameter w such that $\theta = (\lambda, w)$ is not rejected



Profile projections: computing

Idea: compute $\min_{w|\lambda}$ for each candidate target parameter value & compare to threshold
 → if max within threshold then exists a pair in confidence set (why?)

$$\text{Profile}(\lambda) = \min_{w|\lambda} \left\{ \frac{1}{\sigma^2} \|y^{\text{obs}} - y(w, \lambda)\|^2 \right\}$$

Can generalise to profile likelihood functions & other fit measures



Example

```
# plot profile paths
nuisance_K = [K_grid[np.argmax(row)] for row in nllikes]
nuisance_r = [r_grid[np.argmax(col)] for col in nllikes.T]

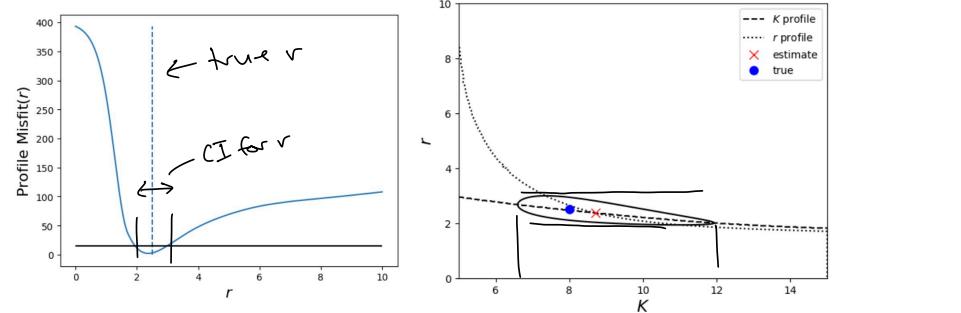
contours = plt.contour(K_grid, r_grid, nllikes, levels=[delta], colors='black')

plt.xlabel(r'$SKS$', fontsize=16)
plt.ylabel(r'$sr$', fontsize=16)

plt.plot(K_grid,nuisance_r,label='$KS$ profile',color='k',linestyle='--')
plt.plot(nuisance_K[2:]:r_grid[2:],label='$r$ profile',color='k',linestyle=':')

plt.plot(prm_est[1],prm_est[0],markersize=8,color='r',marker='x',linestyle='None',label='estimate')
plt.plot(prm_true[1],prm_true[0],markersize=8,color='b',marker='o',linestyle='None',label='true')

plt.legend()
plt.show()
/ 0.1s
```



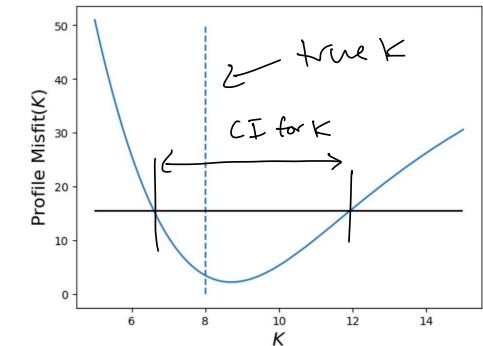
$$CI(r) : [1.94, 2.99]$$

$$\text{true } r : 2.5$$

```
# project via 'manually' profiling over grid cf resolving optimisation.
misfit_K = np.min(nllikes, axis=0)
plt.plot(K_grid, misfit_K)

plt.xlabel(r'$SKS$', fontsize=16)
plt.ylabel('Profile Misfit(K)', fontsize=16)
plt.vlines(x=k_true, ymin=0, ymax=np.max(fit_K),linestyles='--')
plt.hlines(y=delta,xmin=k_min,xmax=k_max,linestyles='-',color='k')
K_llim = K_grid[misfit_K <= delta][0]
K_ullim = K_grid[misfit_K <= delta][-1]
print((K_llim,K_ullim))
plt.show()
/ 0.1s

np.float64(6.603206412825651), np.float64(11.93386773547894))
```



$$CI(K) : [6.60, 11.93]$$

$$\text{true } K : 8$$

Unknown error?

If the error is unknown, we can (try to) estimate as another parameter

→ OK for small dim problems!

But need to think about bias & regularisation for very flexible functions

↳ essentially our regularisation parameter

→ need either another est. eqⁿ for σ or full likelihood function

↳ can't min $\frac{1}{\sigma^2} \|y - f(\theta)\|^2$ in both θ & σ !

→ another approach is to estimate via different empirical model then treat as fixed in sci model

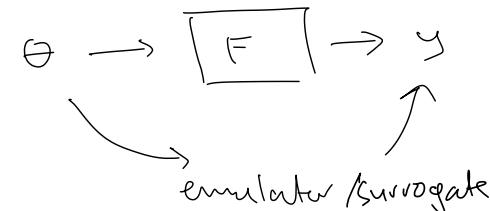
↳ underestimates overall uncert but often 'OK'

Expensive models

For expensive models we can't afford to run as many times and/or compute derivatives

Strategies include:

construct 'emulator' / 'surrogate' from limited budget of runs



- e.g.
- Gaussian process
 - neural net
 - spline

Pay cost up front but can be cheaper long term as emulator cheap to evaluate

+ 'amortised'

stochastic simulation models

$$\theta \rightarrow [F] \rightarrow Y \sim P(Y; \theta)$$

Don't know $P(Y; \theta)$ in closed form but can simulate samples

Options

- Can often find approximate loss function with distribution independent of parameter
 - ↳ compute dist. for some est.
 - ↳ invert hypothesis test
- Construct 'surrogate' $P(Y; \theta)$
 - eg 'synthetic likelihood'
- simulation-based: ABC, LF2I etc

See Canvas

Scientific machine learning

Hybrid { flexible data model
scientific regularisation
↳ not exactly satisfied

→ fg assignment:

spline fit + ODE penalty

Current hype: 'PINNs'

Physics-informed neural networks

→ We'll look at neural nets in last lecture